

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
import pandas as pd
```

```
price = pd.read_csv("/content/drive/MyDrive/Datasets/housing.csv")
```

```
price.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
price.dropna(inplace=True)
```

```
print(price)
```

```
longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0      -122.23    37.88              41.0         880.0         129.0
1      -122.22    37.86              21.0        7099.0        1106.0
2      -122.24    37.85              52.0        1467.0         190.0
3      -122.25    37.85              52.0        1274.0         235.0
4      -122.25    37.85              52.0        1627.0         280.0
...      ...      ...      ...      ...      ...
20635   -121.09    39.48              25.0        1665.0         374.0
20636   -121.21    39.49              18.0         697.0         150.0
20637   -121.22    39.43              17.0        2254.0         485.0
20638   -121.32    39.43              18.0        1860.0         409.0
20639   -121.24    39.37              16.0        2785.0         616.0

population  households  median_income  median_house_value  \
0         322.0        126.0          8.3252        452600.0
1        2401.0       1138.0          8.3014        358500.0
2         496.0        177.0          7.2574        352100.0
3         558.0        219.0          5.6431        341300.0
4         565.0        259.0          3.8462        342200.0
...      ...      ...      ...      ...
20635        845.0       330.0          1.5603         78100.0
20636        356.0       114.0          2.5568         77100.0
20637       1007.0       433.0          1.7000         92300.0
20638        741.0       349.0          1.8672         84700.0
20639       1387.0       530.0          2.3886         89400.0

ocean_proximity
0      NEAR BAY
1      NEAR BAY
2      NEAR BAY
3      NEAR BAY
4      NEAR BAY
...      ...
20635      INLAND
20636      INLAND
20637      INLAND
20638      INLAND
20639      INLAND
```


```
[20433 rows x 10 columns]
```

```
from sklearn.model_selection import train_test_split
```

```
X = price.drop(['median_house_value'], axis=1)
y = price['median_house_value']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
train_data = X_train.join(y_train)
test_data = X_test.join(y_test)
train_data
```



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population
4355	-118.39	34.09	41.0	730.0	126.0	230.0
4209	-118.25	34.11	52.0	125.0	42.0	99.0
17342	-120.40	34.85	26.0	2384.0	385.0	1323.0
4199	-118.24	34.12	34.0	80.0	26.0	125.0
2445	-119.64	36.56	34.0	576.0	117.0	363.0
...
19053	-121.69	38.16	33.0	1808.0	363.0	824.0
20631	-121.40	39.33	15.0	2655.0	493.0	1200.0
18626	-121.95	37.11	21.0	2387.0	357.0	913.0
4377	-118.27	34.09	48.0	1527.0	295.0	589.0
6698	-118.10	34.14	26.0	6262.0	1645.0	3001.0


16346 rows × 10 columns

Next steps:

[Generate code with train_data](#)

[View recommended plots](#)

```
train_data.ocean_proximity.value_counts()
```




```
ocean_proximity
<1H OCEAN    7225
INLAND        5180
NEAR OCEAN    2097
NEAR BAY      1841
ISLAND         3
Name: count, dtype: int64
```

```
train_data = train_data.join(pd.get_dummies(train_data.ocean_proximity, dtype=int)).drop(['ocean_proximity'], axis=1)
test_data = test_data.join(pd.get_dummies(test_data.ocean_proximity, dtype=int)).drop(['ocean_proximity'], axis=1)
```

```
X_train, y_train = train_data.drop(['median_house_value'], axis=1), train_data['median_house_value']
X_test, y_test = test_data.drop(['median_house_value'], axis=1), test_data['median_house_value']
```


```
from sklearn.ensemble import RandomForestRegressor
```

```
forest = RandomForestRegressor()
forest.fit(X_train, y_train)
```




```
▼ RandomForestRegressor
RandomForestRegressor()
```

```
forest.score(X_test, y_test)
```



```
0.8228961965688331
```

```
!pip install lime
```



```
Requirement already satisfied: lime in /usr/local/lib/python3.10/dist-packages (0.2.0.1)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from lime) (3.7.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from lime) (1.25.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from lime) (1.11.4)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from lime) (4.66.4)
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.10/dist-packages (from lime) (1.2.2)
```

Requirement already satisfied: scikit-image>=0.12 in /usr/local/lib/python3.10/dist-packages (from lime) (0.19.3)
Requirement already satisfied: networkx>=2.2 in /usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (3.3)
Requirement already satisfied: pillow!=7.1.0,!>=7.1.1,!>=8.3.0,>=6.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (2.31.6)
Requirement already satisfied: imageio>=2.4.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (2.31.6)
Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (2024.5.10)
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (1.6.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from scikit-image>=0.12->lime) (24.0)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18->lime) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18->lime) (3.5.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.5.0->lime) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.5.0->lime) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.5.0->lime) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.5.0->lime) (1.4.5)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.5.0->lime) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.5.0->lime) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib>=3.5.0->lime) (1.16.0)

```
import lime
import lime.lime_tabular

X_train_np = X_train.to_numpy().astype(float)

explainer = lime.lime_tabular.LimeTabularExplainer(X_train_np,
    feature_names=X_train.columns,
    class_names=["median_house_value"],
    mode='regression'
)

exp = explainer.explain_instance(X_test.iloc[0], forest.predict)

exp.show_in_notebook(show_table=True)
```

